



# Motion Transfer Pipeline Overview

Image-to-video (i2v) motion transfer for Wan T12V-5B: learn motion from a **driver video** (generated from `cat.png`), then generate a new video from `corgi.png` that follows the same motion, guided by text.

All tasks remain **i2v**: the **first-frame latent is fixed** from the input image; only later frames are denoised.

---

## Pipeline graph (Mermaid)

```
graph TD
    subgraph TRAIN ["Training (train_motion.py)"]
        DV["Driver video<br/>klings_20260512.mp4"]
        CAT["Ref image<br/>cat.png"]
        PROMPT["Text prompt"]
        DV --> ALIGN_T["Align & resize frames<br/>to cat.png aspect / max_area"]
        CAT --> ALIGN_T
        ALIGN_T --> VAE_ENC_T["Wan VAE encode<br/>full video → target_z"]
        ALIGN_T --> VAE_1_T["Wan VAE encode<br/>frame 0 → cond_z"]
        VAE_ENC_T --> CLIP_ENC_T["MotionClipEncoder<br/>OpenCLIP ViT + temporal TF"]
        CLIP_ENC_T --> MT_T["32 motion tokens<br/>512-d"]
        MT_T --> ADAPTER_T["MotionContextAdapter<br/>MLP → 3072-d"]
        PROMPT --> T5_T["T5-XXL text encoder"]
        T5_T --> CTX_T["Text context tokens"]
        VAE_1_T --> FM_T["Flow-matching loss<br/>noisy latent + mask"]
        ADAPTER_T --> DIT_T["WanModelWithMotion (frozen)<br/>motion tokens || text → cross-attn"]
        CTX_T --> DIT_T
        FM_T --> DIT_T
        DIT_T --> LOSS["MSE(pred, noise - target_z)"]
        LOSS --> UPD["Update motion encoder + adapter only"]
        UPD --> CKPT["checkpoints/motion_transfer.pt"]
    end

    subgraph INFER ["Inference (generate_motion.py)"]
        DV2["Driver video"]
        CAT2["motion_ref_image<br/>cat.png (crop align)"]
        CORGI["Target image<br/>corgi.png"]
        PROMPT2["Text prompt"]
        CKPT2["motion_transfer.pt"]
        DV2 --> ALIGN_I["Align frames to cat.png geometry"]
        CAT2 --> ALIGN_I
        ALIGN_I --> CLIP_ENC_I["MotionClipEncoder"]
        CKPT2 --> CLIP_ENC_I
        CLIP_ENC_I --> ADAPTER_I["MotionContextAdapter"]
        CLIP_ENC_I --> MT_I["Motion tokens"]
        ADAPTER_I --> CORGI
        CORGI --> PRE_I["Resize / center-crop"]
    end
```

```
PRE_I --> VAE_1_I["VAE encode frame 0 → z"] PRE_I --> I2V["Wan T12V i2v sampling loop"] PROMPT2 --> T5_I["T5 encoder"] T5_I --> CTX_I["Text context"] ADAPTER_I --> DIT_I["WanModelWithMotion"] CTX_I --> DIT_I VAE_1_I --> I2V MT_I --> DIT_I I2V --> DIT_I DIT_I --> VAE_DEC["VAE decode"] VAE_DEC --> OUT["Output video<br/>corgi + transferred motion"] end subgraph ARCH["Motion module detail"] F1["Frame t"] --> CLIP_V["CLIP ViT-B/32<br/>frozen"] F2["Frame t−1"] --> CLIP_V CLIP_V --> DELTA["Δ = feat_t − feat_{t−1}"] CLIP_V --> CONCAT["concat feat, Δ"] DELTA --> CONCAT CONCAT --> TEMP["Temporal Transformer"] TEMP --> Q["32 learnable queries"] Q --> CROSS["Cross-attention"] TEMP --> CROSS CROSS --> TOK["32 × 512 motion tokens"] TOK --> MLP["MLP × scale → 3072"] MLP --> INJECT["Prepend to T5 context<br/>in every DiT block cross-attn"] end
```

---

# High-level goal

| Asset                | Path                           | Role                                                  |
|----------------------|--------------------------------|-------------------------------------------------------|
| Driver video         | examples/kling_20260512.mp4    | Source of motion (and training reconstruction target) |
| Original first frame | examples/cat.png               | Spatial alignment reference for driver / training     |
| New appearance       | examples/corgi.png             | First frame at inference                              |
| Checkpoint           | checkpoints/motion_transfer.pt | Trained motion encoder + adapter                      |

---

# Phase A — Training ( `train_motion.py` )

## Inputs

| Input                                                | Role                                             |
|------------------------------------------------------|--------------------------------------------------|
| Driver video<br>( <code>klings_20260512.mp4</code> ) | Ground-truth motion + appearance to reconstruct  |
| Ref image ( <code>cat.png</code> )                   | Defines crop / aspect (same as original i2v)     |
| Text prompt                                          | Same semantic description as original generation |
| Wan T12V-5B checkpoint                               | Frozen backbone (VAE, T5, DiT)                   |

## Step 1 — Video preprocessing

- Sample `T` frames (e.g. 49 with `--low_vram`, 81 at full settings).
- Resize and center-crop each frame to match `cat.png` layout and `max_area` (e.g. `832×480` or `1280×704` ).
- Tensor layout: `[B, T, 3, H, W]` in `[-1, 1]` .

## Step 2 — Latent targets (frozen Wan VAE)

- Encode full video → `target_z` , shape `[C_z, T_z, H_z, W_z]` (e.g. `48×13×44×34` ).
- Encode first frame only → `cond_z` (i2v conditioning).
- Optional cache: `checkpoints/motion_transfer_latents.pt` (skips re-encoding on reruns).

## Step 3 — Motion features (trainable `MotionClipEncoder` )

1. Per frame: OpenCLIP ViT-B/32 → 512-d embedding (CLIP weights frozen).
2. Temporal delta: `feat_t - feat_{t-1}` .
3. Concatenate `[feat, delta]` → linear projection → 512-d per timestep.
4. Small temporal Transformer over time.

5. 32 learnable query tokens + cross-attention → **32 motion tokens** × **512**.

## Step 4 — Text conditioning (frozen T5)

- Prompt → T5-XXL → text tokens [L, 4096] → Wan text embedding → [L, 3072] .

## Step 5 — Inject motion into DiT ( WanModelWithMotion )

- MotionContextAdapter : MLP maps 512 → 3072 (Wan hidden size), with a learnable scale.
- **Concatenate** motion tokens **before** text tokens in cross-attention context (same idea as Wan-Animate CLIP tokens).
- Wan DiT (5B, 30 layers) remains **frozen**.

## Step 6 — Flow-matching training step

- Sample noise and timestep  $t$  .
- Build noisy latent: interpolate `target_z` and noise; **overwrite first latent frame** with `cond_z` (i2v mask).
- DiT predicts velocity / noise residual.
- Loss:  $\text{MSE}(\text{prediction}, \text{noise} - \text{target\_z})$  .
- Backprop **only** into motion encoder (optional) + motion adapter.
- Gradient checkpointing on DiT blocks during training to reduce VRAM.

## Training output

- checkpoints/motion\_transfer.pt — weights for MotionClipEncoder + MotionContextAdapter .

# Example command

```
python train_motion.py \
  --ckpt_dir ../Wan2.2-TI2V-5B \
  --driver_video examples/kling_20260512.mp4 \
  --ref_image examples/cat.png \
  --prompt "the cute cat raises its right paw and licking the paw repeatedly, satisfied and" \
  --frame_num 81 \
  --steps 500 \
  --low_vram \
  --output checkpoints/motion_transfer.pt
```

## Phase B — Inference ( generate\_motion.py )

### Inputs

| Input                       | Role                                  |
|-----------------------------|---------------------------------------|
| corgi.png                   | New appearance (first frame)          |
| cat.png as motion_ref_image | Same crop geometry as driver video    |
| Driver video                | Motion source (or cached motion .pt ) |
| motion_transfer.pt          | Learned motion modules                |
| Prompt                      | Describes corgi action                |

### Steps

1. Extract motion from driver video (aligned with cat.png ) → 32 motion tokens.
2. Load corgi.png → preprocess → VAE encode first frame → z .
3. T5-encode prompt.
4. Standard Wan i2v diffusion loop (UniPC / flow matching):
  - Latent noise for future frames; frame 0 fixed from z .
  - Each step: DiT with **motion tokens + text** (CFG: motion disabled on

unconditional branch).

5. VAE decode → output MP4.

## Design choice

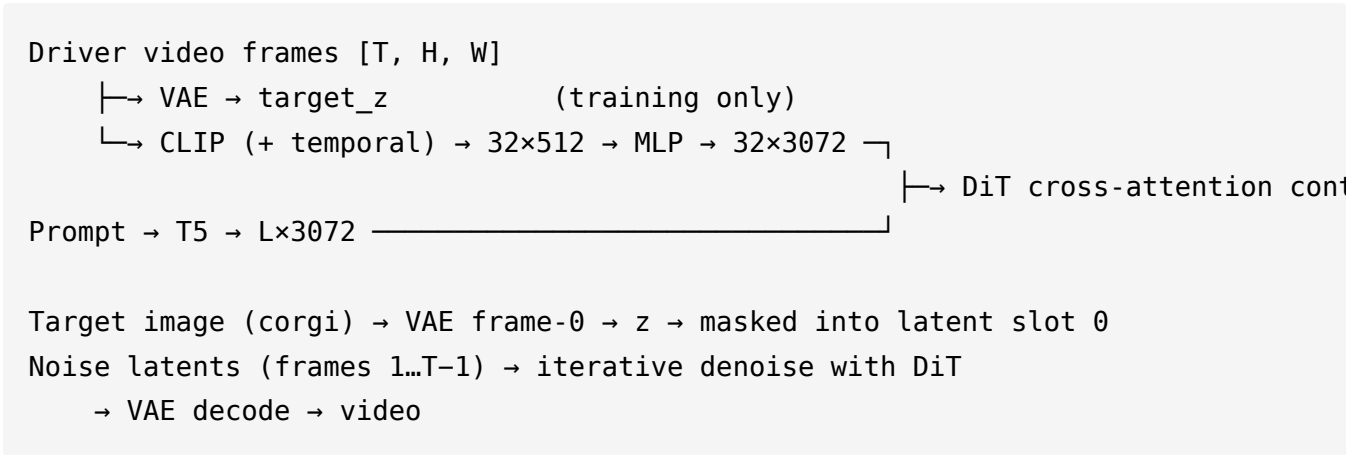
- **Motion** is largely appearance-agnostic in CLIP space.
- **Appearance** comes from `corgi.png` + prompt.
- **Motion** comes from driver video features.

## Example command

```
python generate_motion.py \  
  --ckpt_dir ../Wan2.2-TI2V-5B \  
  --motion_ckpt checkpoints/motion_transfer.pt \  
  --image examples/corgi.png \  
  --motion_ref_image examples/cat.png \  
  --motion_video examples/kling_20260512.mp4 \  
  --prompt "the cute corgi raises its right paw and licking the paw repeatedly, satisfied" \  
  --size 1280*704 \  
  --frame_num 81 \  
  --offload_model True \  
  --convert_model_dtype \  
  --t5_cpu \  
  --save_file output/corgi_motion.mp4
```

---

# Tensor-level data flow



## Module map (code)

| Component              | File                                        |
|------------------------|---------------------------------------------|
| MotionClipEncoder      | wan/modules/motion_transfer/clip_encoder.py |
| MotionContextAdapter   | wan/modules/motion_transfer/adapter.py      |
| WanModelWithMotion     | wan/modules/motion_transfer/model.py        |
| WanTI2VMotion pipeline | wan/textimage2video_motion.py               |
| Training script        | train_motion.py                             |
| Inference script       | generate_motion.py                          |
| Motion cache utility   | extract_motion.py                           |

# VRAM and resolution

| Mode                                | Resolution | Frames | GPU notes                                                 |
|-------------------------------------|------------|--------|-----------------------------------------------------------|
| Training<br><code>--low_vram</code> | 832×480    | 49     | Fits ~24GB; DiT offloaded between steps                   |
| Training (full)                     | 1280×704   | 81     | Needs 40GB+; OOM on 24GB                                  |
| Inference                           | 1280×704   | 81     | Same as base Wan i2v with<br><code>--offload_model</code> |

Training and inference resolutions may differ; the motion adapter maps tokens into DiT context and is not tied to a single latent grid size.

---

## Suggested figure layout (for external diagram tools)

### Panel 1 — Training

- **Left:** driver video + cat reference → preprocessing.
- **Center:** three branches (VAE latents, CLIP motion, T5 text) merging into frozen DiT.
- **Right:** loss arrow into small trainable block (Motion Encoder + Adapter).

### Panel 2 — Inference

- **Top:** driver → motion branch.
- **Bottom:** corgi → appearance branch.
- **Merge** at DiT → VAE decode → output video.

### Panel 3 — Motion encoder (zoom-in)

Frames → CLIP → temporal deltas → Temporal Transformer → 32 query tokens → cross-attention → MLP → prepend to text context in every DiT block.



---

## Related docs

- Quick start: `motion_transfer.md`
- Base Wan i2v: `inference.md`